

EE496 Homework 1 - Training Multi-Layer Perceptron

Berken Utku Demirel - 2166221

March 20, 2020

1 Basic Concepts

1.1 Which Function ?

What function does the MLP classifier of Scikit-Learn [2] approximates ?

Multi-layer Perceptron (MLP) is a supervised learning algorithm that approximates the function $f(\cdot) : R^m \rightarrow R^o$ by training on a given dataset, where m and o are the dimensions of input and output respectively.

How is the loss defined to approximate that function ?

Multi-layer perceptron(MLP) uses Cross-Entropy as a loss function for classification. In binary classification, cross-entropy can be calculated as in equation 1 where the \hat{y} is the predicted probability observation and y is the binary indicator if class label is the correct classification for that observation.

$$Loss(\hat{y}, y, W) = -y \ln \hat{y} - (1 - y) \ln(1 - \hat{y}) \quad (1)$$

If the classification is multiclass, which is in our case, the loss for each class is calculated and sum the result as in equation 2 where the M is the number of classes.

$$Loss = - \sum_{k=1}^M y_k \ln \hat{y}_k \quad (2)$$

The reason why there is a defined loss to approximate that function is that the program learn by means of a loss function. It's a method of evaluating how well specific algorithm models the given dataset. If predictions deviates too much from actual results, loss function would result in a very large number. Gradually, with the help of some optimization function, loss function learns to reduce the error in prediction. Furthermore, an important aspect of cross-entropy loss is that it penalizes the predictions that are confident but wrong since when actual label is 1 ($y(i) = 1$), second half of function disappears whereas in case actual label is 0 ($y(i) = 0$) first half is dropped off.

1.2 Gradient Computation

The equation that describes the update rule for weights is given in equation 3,

$$w_{k+1} = w_k - \gamma \nabla_w \mathcal{L}|_{w=w_k} \quad (3)$$

If we solve this equation for the gradient of the loss, we would end up with the equation 4.

$$\nabla_w \mathcal{L}|_{w=w_k} = \frac{w_k - w_{k+1}}{\gamma} \quad (4)$$

1.3 Some Training Parameters and Basic Parameter Calculations

1.3.1

Batch is a hyperparameter that defines the number of samples to work through before updating the internal model parameters.

Epoch is also a hyperparameter that defines the how many times the learning algorithm work through entire training dataset.

1.3.2

If the dataset has N samples and the batch size is B , the number of batches per epoch is N/B .

1.3.3

Stochastic gradient decent (SGD) approximate the gradient using only one data point. So, SGD iterates only once for a batch. The number of SGD iterations in an epoch is $\text{floor}(N/B)$. Therefore, the total number of SGD iteration in the training is $E \times \text{floor}(N/B)$.

1.4 Computing Number of Parameters of an MLP Classifier

To calculate the number of parameters in an MLP classifier, we have to find the total number of the weights and biases of the neural network. The calculation of the number of weights, which corresponds to the total number of connections between the layers, is given in equation 5.

$$\# \text{ of weights} = D_{in} \times H_1 + D_{out} \times H_k + \sum_{k=2}^{k-1} H_k \times H_{k+1} \quad (5)$$

Furthermore, the biases in each layer should be computed as in equation 6.

$$\# \text{ of biases} = \sum_{k=1}^k H_k \quad (6)$$

The sum of these two equations yields the number of parameters of the MLP, which is given in equation 7.

$$\# \text{ of parameters of MLP} = D_{in} \times H_1 + D_{out} \times H_k + \sum_{k=2}^{k-1} H_k \times H_{k+1} + \sum_{k=1}^k H_k \quad (7)$$

2 Experimenting MLP Architectures

2.1 Experimental Work

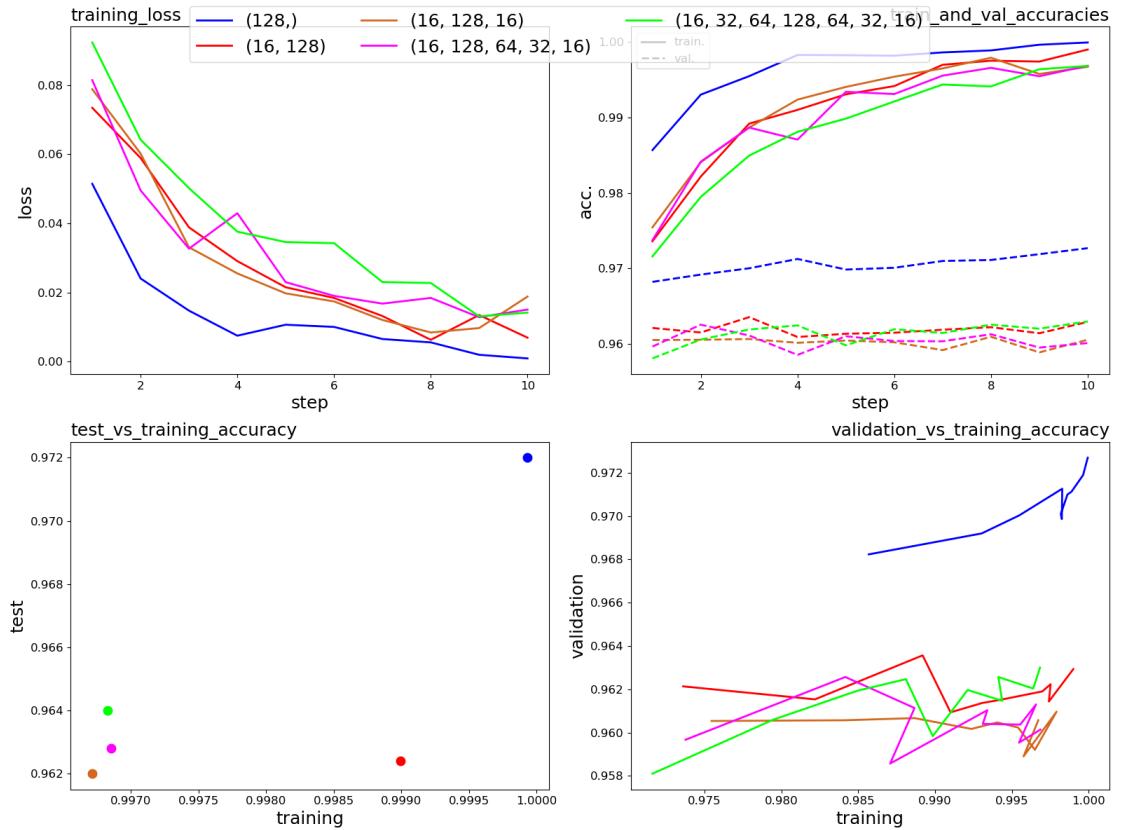


Figure 1: The result of *part2Plot* function with 100 epoch

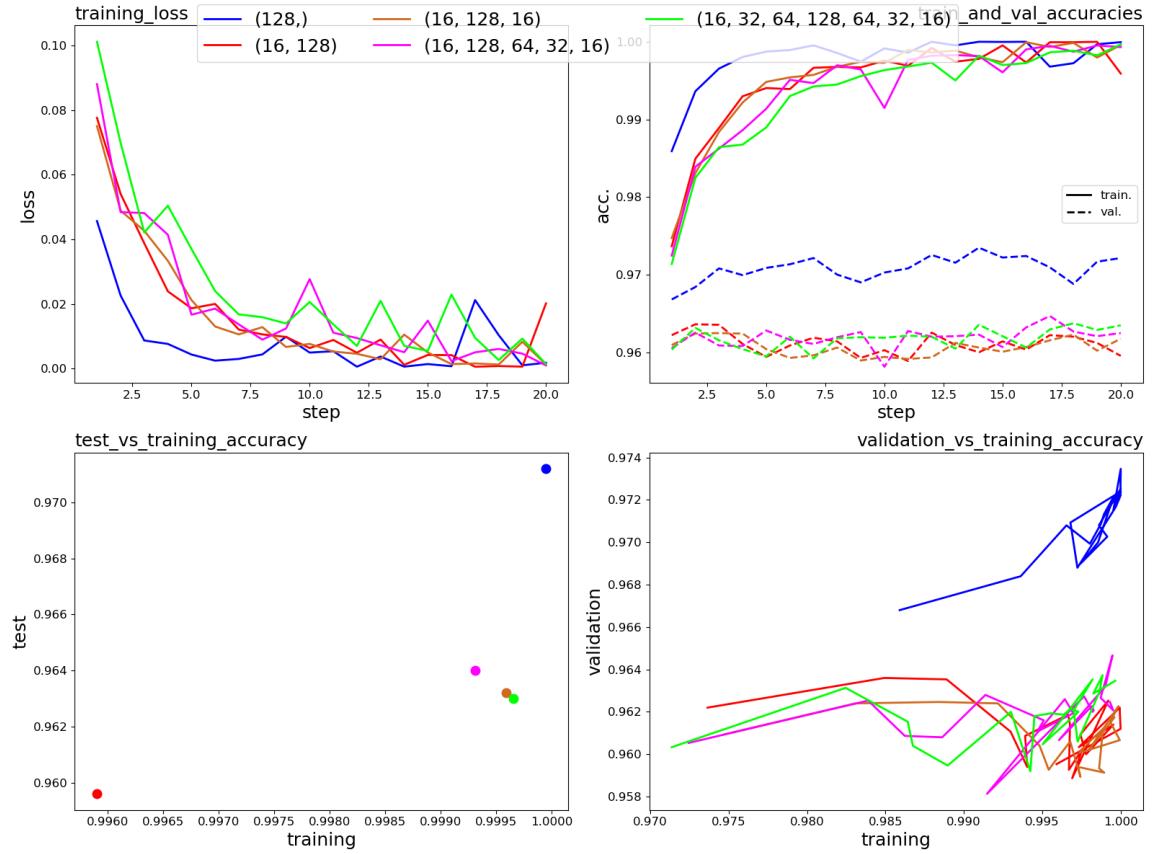


Figure 2: The result of the *part2Plot* function with 200 epoch

The figure 2 shows that the result of the part2 when the MLP has trained for 200 epochs. This plot is not wanted in the homework, however, in order to compare the generalization performance of the classifiers, it is obtained.

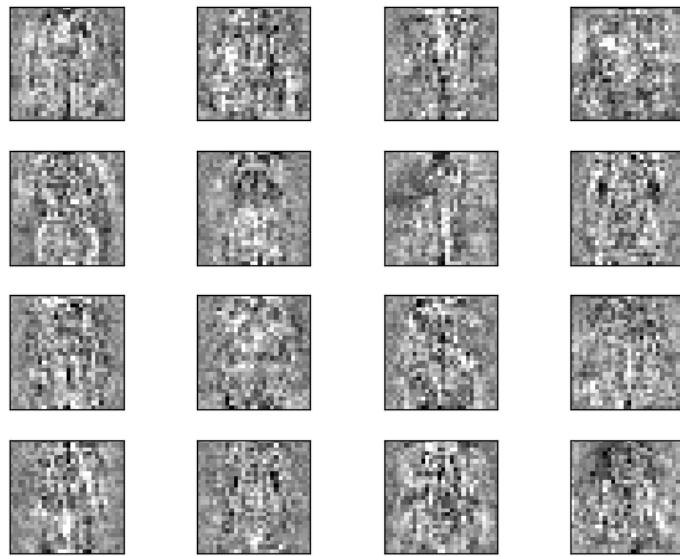


Figure 3: The result of the *visualizeWeights* function for the arch_2

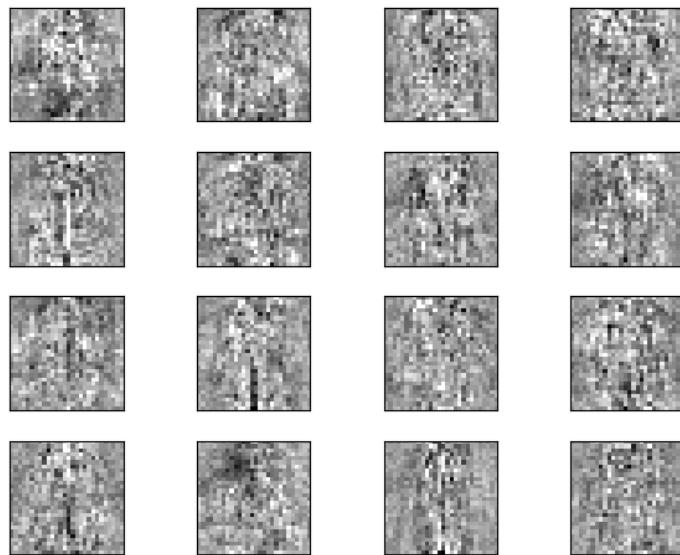


Figure 4: The result of the *visualizeWeights* function for the arch_3



Figure 5: The result of the *visualizeWeights* function for the arch_5



Figure 6: The result of the *visualizeWeights* function for the arch_7

2.2 Discussions

2.2.1 What is the generalization performance of a classifier?

Generalization performance of a classifier refers to your model's ability to adapt properly to new, previously unseen data, drawn from the same distribution as the one used to create the model.

2.2.2 Which plots are informative to inspect generalization performance?

We can make an inference about the generalization of the trained neural network by looking at the validation and test accuracies curve. Since the data, which is used for validation and test, are previously unseen data, drawn from the same distribution as the one used to create the model.

2.2.3 Compare the generalization performance of the architectures

When we look at the validation and test accuracies curve, we can observe that the 'arch_1' reaches the highest values in both. So the generalization performance of the 'arch_1' is superior to others. The other architectures are not so different in terms of the generalization. However, the 'arch_7', which is the deepest one, shows the lowest generalization performance amongst the others which can be seen from the figure 2. When the number of the epoch is increased to 200, the training accuracy of the 'arch_7' is better than 'arch_1' whereas the validation accuracy of the 'arch_7' is much worse than the 'arch_1'.

2.2.4 How does the number of parameters affect the classification and generalization performance?

When a neural network has more parameters, it has sufficient capacity to overfit the training data. Reducing the number of parameters in the network will increase the generalization capacity of the network. In general, with a fixed number of training patterns, overfitting can occur when the model has too many parameters (too many degrees of freedom).

2.2.5 How does the depth of the architecture affect the classification and generalization performance?

More layers offer more opportunity for hierarchical re-composition of abstract features learned from the data. However, the deep architectures need more data for training. Also, when we increase the size of the network, we'll introduce more parameters that network needs to learn, and hence increasing the chances of overfitting.

2.2.6 Considering the visualizations of the weights, are they interpretable?

The weights are similar to the inputs, which are fed to the network. However, when we look at the visualizations of the weights for the deeper networks especially 6, some of the figures are not detailed, which means that the network couldn't learn the features entirely for the classification.

2.2.7 Can you say whether the units are specialized to specific classes

2.2.8 Weights of which architecture are more interpretable?

When we look at the figures from 3 to 6, we can deduce that some of the images in the figure 6 are similar to dress and footwear input images.

2.2.9 Considering the architectures, comment on the structures (how they are designed). Can you say that some architecture are akin to each other? Compare the performance of similarly structured architectures and architectures with different structure.

The hidden layer size of all architectures is chosen as the power of 2. The most shallow one is composed of 1 hidden layer, whereas the deepest one is designed as consisting of 7 hidden layers. So we can separate these 5 different architectures as the shallow and deep neural network. When we compare the figures 1 and 2, we can easily observe that when the neural network is going to deeper, the new parameters are introduced to that network, and it is more prone to overfitting, thus the generalization performance of these neural networks is worse than the shallow one.

2.2.10 Which architecture would you pick for this classification task? Why?

I would choose the 'arch_1' because it can be observed from figures 1 and 2, the generalization performance of the classifier is the best amongst the others.

3 Experimenting Activation Functions

3.1 Experimental Work

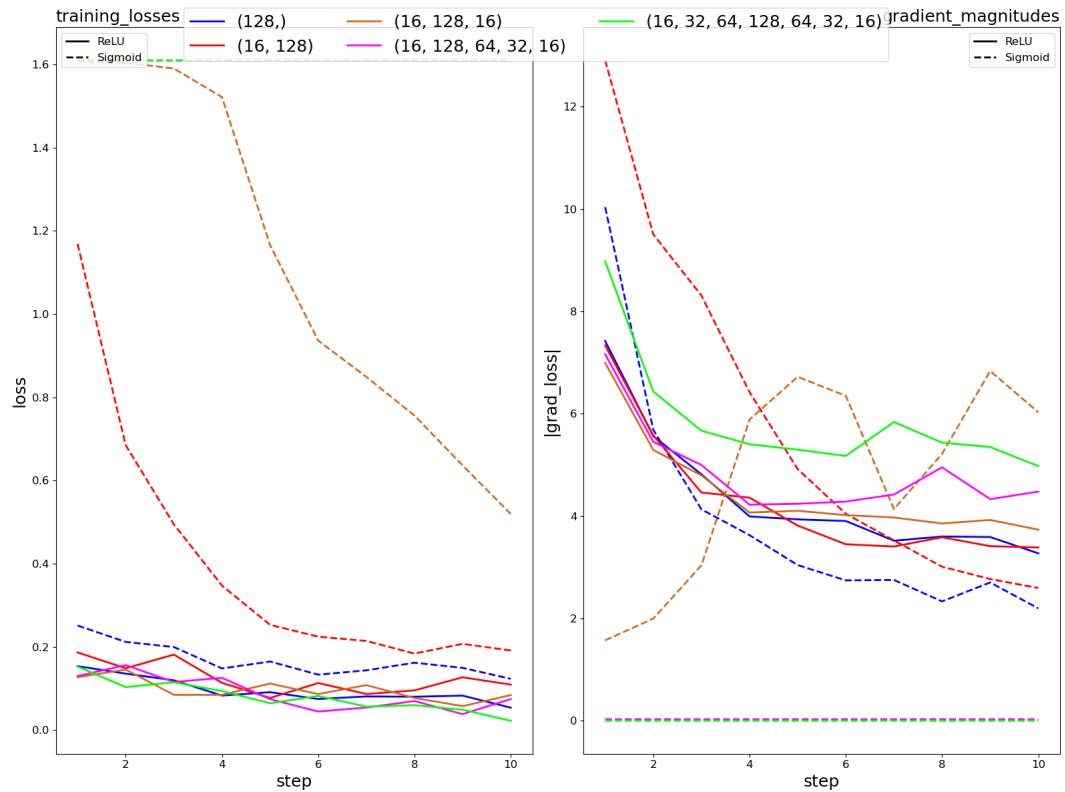


Figure 7: The result of *part3Plot* function

3.2 Discussions

3.2.1 How is the gradient behavior in different architectures? What happens when depth increases?

When we look at the figure 7, we can observe that when the network architecture is deeper, the magnitude of the loss gradient is decreasing especially for the architectures which use sigmoid function. As more layers using certain activation functions are added to neural networks, the gradients of the loss function approaches zero, making the network hard to train.

3.2.2 Why do you think that happens?

Certain activation functions, like the sigmoid function, squishes a large input space into a small input space between 0 and 1. Therefore, a significant change in the input of the sigmoid function will cause a minimal change in the output. Hence, the derivative becomes small. This is not a big problem for shallow networks like 'arch_1'. However, when more layers are used, it can cause the gradient to be too small for training because n hidden layers use the sigmoid function, n small derivatives are multiplied together. Thus, the gradient decreases exponentially as we propagate down to the initial layers.

3.2.3 What might happen if we do not scale the inputs to the range [-1.0,1.0]?

If we don't scale the data, then the convergence will be slower. The training time will be more compared to training using normalized data. Also, the gradient descent will take a longer time to converge compared to train by using scaled data. In addition, standardizing the inputs can reduce the chances of getting stuck in local optima.

4 Experimenting Learning Rate

4.1 Experimental Work

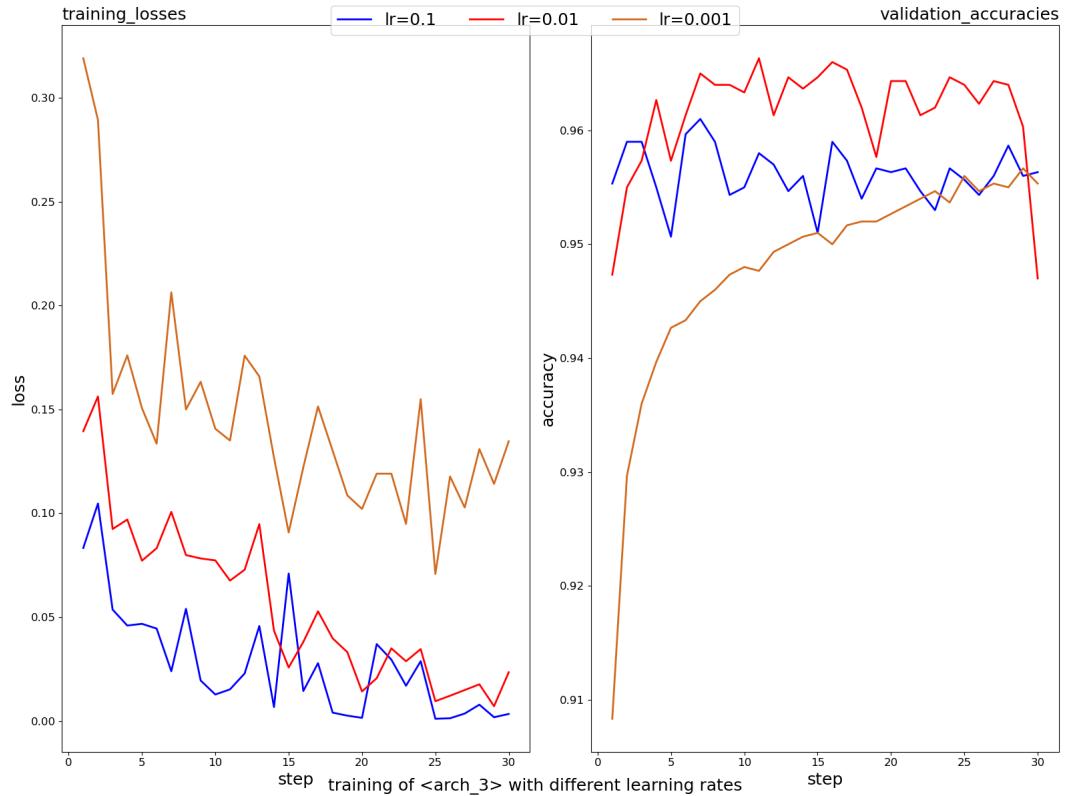


Figure 8: The result of *part4Plot* function for 'arch_3'

Figure 8 shows the result of the *part4Plot* function for the three different MLP classifier with learning rates 0.1, 0.01 and 0.001, respectively.

Figure 9 represents the validation accuracy when we set the learning rate to 0.01 and continue training until 200 epochs.

Figure 10 represents the validation accuracy when the learning rate is changed to twice first 0.01 then 0.001 and continue training until 200 epochs.

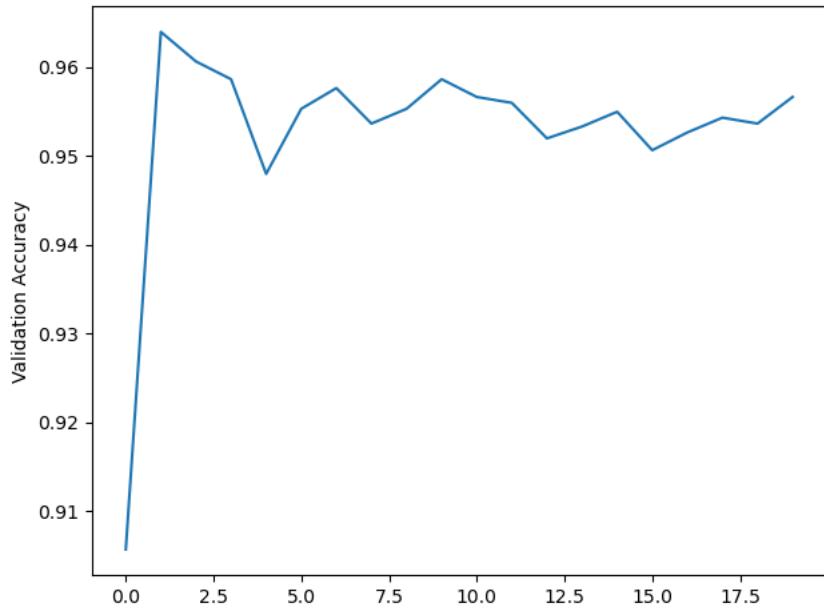


Figure 9: The validation curve for the part 4.3 with $\gamma = 0.01$

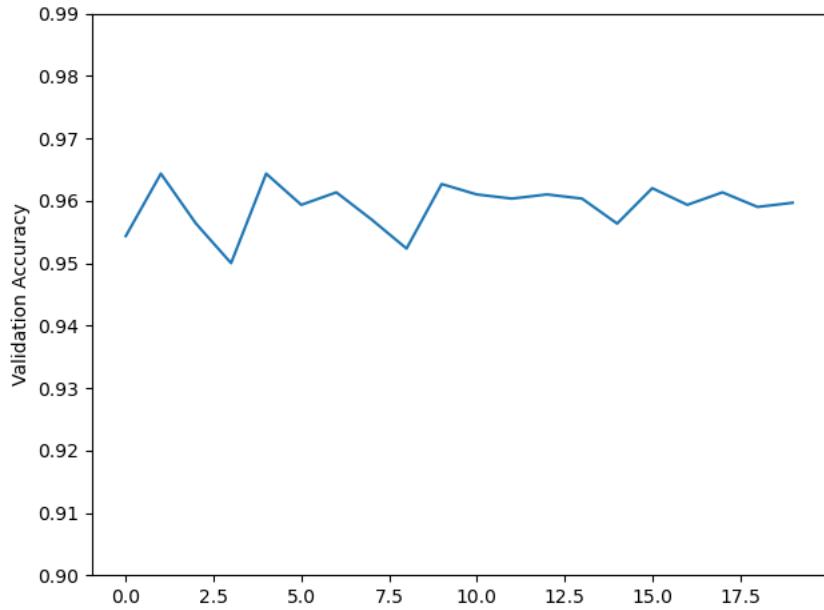


Figure 10: The validation curve for the part 4.6 with $\gamma = 0.001$

4.2 Discussions

4.2.1 How does the learning rate affect the convergence speed?

The learning rate controls how much to change the model in response to the estimated error each time the model weights are updated. Choosing the learning rate is important since if it is too small, the learning results in a long training process that could get stuck. In contrast, a large learning rate might result in learning a sub-optimal set of weights too fast or an unstable training process. The figure 8 shows that the architecture, which has a bigger learning rate, converges to its maximum validation accuracy in a smaller step.

4.2.2 How does the learning rate affect the convergence to a better point?

When the learning rate is too large, the learning may end in a not-optimal point since as the learning rate increases, the network becomes more unstable. This fact can also be seen in figure 8. The validation accuracy of the architecture with $\gamma = 0.1$ is behaving unstable(lots of peaks), whereas the architecture, which has a learning rate of 0.001, is more stable.

4.2.3 Does your scheduled learning rate method work? In what sense?

Yes, It works. When we look at the figures 9 and 10, we can easily see that the instability of the model is getting smaller in the figure 10 that the scheduled learning rate is applied. Also, the validation accuracy(0.96) is better than the model in which there is no scheduled learning rate.

4.2.4 Compare the accuracy and convergence performance of your scheduled learning rate method with ADAM.

When the ADAM method was used, the best accuracy is 0.9723, which can be seen from figure 1. The accuracy of the scheduled learning rate method is 0.9933. So, we can deduce that the scheduled learning rate method is better than ADAM in terms of accuracy. For the comparison of the convergence performance, both architectures work well since, at the end of the steps, they are stable. However, when we have used the scheduled learning rate the number of epoch which is needed for the convergence is bigger than the ADAM. Therefore, we can make an inference that the convergence performance of the ADAM is better than the scheduled learning rate method.

5 Appendix

```

1 from sklearn.preprocessing import MinMaxScaler
2 from sklearn.model_selection import train_test_split
3 import matplotlib.pyplot as plt
4 from sklearn.neural_network import MLPClassifier
5 import numpy as np
6
7 # Load the train and test images with labels
8 train_images = np.load('train_images.npy')
9 train_labels = np.load('train_labels.npy')
10 test_images = np.load('test_images.npy')
11 test_labels = np.load('test_labels.npy')
12
13 train_images_reshaped = np.reshape(train_images, (30000,
14 28, 28))
14 test_images_reshaped = np.reshape(test_images, (
15 test_images.shape[0], 28, 28))
16
16 scaled_images = np.empty(shape=(30000, 28, 28))
17 scaled_images_test = np.empty(shape=(5000, 28, 28))
18
19 # In order to scale pixel values to [-1, 1], The
20 # MinMaxScaler was used with a range [-1,1]
21 counter = 0
21 scaler = MinMaxScaler(feature_range=(-1, 1))
22
23 for x in train_images_reshaped:
24     scaled_images[counter, :, :] = scaler.fit_transform(x)
25     counter += 1
26
27 counter = 0
28 for x in test_images_reshaped:
29     scaled_images_test[counter, :, :] = scaler.
30         fit_transform(x)
31     counter += 1
32
32 # In order to split 10 percent of the training data set to
33 # validation
33 X_train, X_validate, y_train, y_validate =
34     train_test_split(scaled_images,
35
35     train_labels, test_size=0.1,
36
36     random_state=40, stratify=train_labels)
37
37 # Definition of all architectures

```

```

38 arch_1 = (128,)
39 arch_2 = (16, 128,)
40 arch_3 = (16, 128, 16,)
41 arch_5 = (16, 128, 64, 32, 16,)
42 arch_7 = (16, 32, 64, 128, 64, 32, 16,)
43
44 arch = [arch_1, arch_2, arch_3, arch_5, arch_7]
45
46 list_of_dict = []
47
48 # For loop for all architectures
49 for m in arch:
50     average_loss = np.zeros(10)
51     average_valid_accuracy = np.zeros(10)
52     average_training_accuracy = np.zeros(10)
53     overall_score = np.zeros(10)
54     weights_first_layer = []
55     arc_dict = {}
56     counter_for_list = 0
57     # Repeat the all training for 10 times, then take the
58     # average of these
58     for x in range(0, 10):
59         mlp = MLPClassifier(hidden_layer_sizes=m,
60                             activation='relu',
61                             solver='adam', max_iter=1,
62                             shuffle=True)
63
64         # Flatten input data
65         nsamples, nx, ny = X_train.shape
66         d2_train_dataset = X_train.reshape((nsamples, nx*
67             ny))
68
69         nsamples, nx, ny = X_validate.shape
70         d2_validate_dataset = X_validate.reshape((nsamples
71             , nx*ny))
72
73         valid_accuracy = np.empty(shape=(10,))
74         training_accuracy = np.empty(shape=(10,))
75         Loss = np.empty(shape=(10,))
76
77         # For loop to 100 epochs

```

```

77         counter_10_epoch = 0
78         for i in range(1, 101):
79             # Shuffle training set after each epoch
80             random_perm = np.random.permutation(X_train.
81                 shape[0])
82             # Set the mini_batch indexes to 0 after each
83             epoch
84             mini_batch_index = 0
85
86             # Until all data is fed to network, continue
87             # to feed batches.
88             while True:
89                 indices = random_perm[mini_batch_index:
90                     mini_batch_index + 500]
91                 mlp.partial_fit(d2_train_dataset[indices],
92                     y_train[indices], np.unique(y_train))
93                 mini_batch_index += 500
94
95                 # When the batches are bigger than
96                 # dataset, break the epoch
97                 if mini_batch_index >= X_train.shape[0]:
98                     break
99
100
101             # Record the training loss, training accuracy
102             # , validation accuracy for every 10 steps
103             if i % 10 == 0:
104                 valid_accuracy[counter_10_epoch] = mlp.
105                 score(d2_validate_dataset, y_validate)
106                 training_accuracy[counter_10_epoch] = mlp
107                 .score(d2_train_dataset, y_train)
108                 Loss[counter_10_epoch] = mlp.loss_
109                 counter_10_epoch += 1
110
111             average_loss = (average_loss + Loss)
112             average_valid_accuracy = (average_valid_accuracy
113             + valid_accuracy)
114             average_training_accuracy = (
115                 average_training_accuracy + training_accuracy)
116             overall_score[x] = mlp.score(d2_test_dataset,
117                 test_labels)
118             weights_first_layer.append(mlp.coefs_[0])
119
120             # At the final divide the each parameter to 10, in
121             # order to get average values.
122             average_loss = average_loss / 10

```

```
109     average_valid_accuracy = average_valid_accuracy / 10
110     average_training_accuracy = average_training_accuracy
111     / 10
112
113     # Get the index and value of the best test accuracy
114     best_test_accuracy_index = np.argmax(overall_score)
115     best_accuracy = overall_score[
116         best_test_accuracy_index]
117     best_weights = weights_first_layer[
118         best_test_accuracy_index]
119
120     arc_dict['name'] = m
121     arc_dict['loss_curve'] = average_loss
122     arc_dict['train_acc_curve'] =
123         average_training_accuracy
124     arc_dict['val_acc_curve'] = average_valid_accuracy
125     arc_dict['test_acc'] = best_accuracy
126     arc_dict['weights'] = best_weights
127     list_of_dict.append(arc_dict)
128
129 # Visualization Part
130 from utils import part2Plots, visualizeWeights
131 part2Plots(list_of_dict, save_dir='', filename='',
132             show_plot=True)
133
134 a = list_of_dict[1]['weights']
135 visualizeWeights(a, save_dir='', filename='weigths')
136
137 a = list_of_dict[2]['weights']
138 visualizeWeights(a, save_dir='', filename='weigths')
139
140 a = list_of_dict[3]['weights']
141 visualizeWeights(a, save_dir='', filename='weigths')
```

```
1 from sklearn.preprocessing import MinMaxScaler
2 from sklearn.model_selection import train_test_split
3 import matplotlib.pyplot as plt
4 from sklearn.neural_network import MLPClassifier
5 import numpy as np
6 import copy
7
8 train_images = np.load('train_images.npy')
9 train_labels = np.load('train_labels.npy')
10 test_images = np.load('test_images.npy')
11 test_labels = np.load('test_labels.npy')
12
13 train_images_reshaped = np.reshape(train_images, (30000,
14 28, 28))
14 test_images_reshaped = np.reshape(test_images, (
15 test_images.shape[0], 28, 28))
16
16 scaled_images = np.empty(shape = (30000, 28, 28))
17 scaled_images_test = np.empty(shape=(5000, 28, 28))
18 # In order to scale pixel values to [-1, 1], The
18 MinMaxScaler was used.
19 counter = 0
20 scaler = MinMaxScaler(feature_range=(-1, 1))
21
22 for x in train_images_reshaped:
23     scaled_images[counter, :, :] = scaler.fit_transform(x)
24     counter += 1
25
26 counter = 0
27 for x in test_images_reshaped:
28     scaled_images_test[counter, :, :] = scaler.
28 fit_transform(x)
29     counter += 1
30
31 # In order to split 10 percent of the training data set to
31 validation
32 X_train, X_validate, y_train, y_validate =
32 train_test_split(scaled_images,
33
33     train_labels, test_size=0.1,
34
34     random_state=40, stratify=train_labels)
35
36 arch_1 = (128,)
37 arch_2 = (16, 128,)
```

```

38 arch_3 = (16, 128, 16, )
39 arch_5 = (16, 128, 64, 32, 16, )
40 arch_7 = (16, 32, 64, 128, 64, 32, 16, )
41
42 arch = [arch_1, arch_2, arch_3, arch_5, arch_7]
43
44 list_of_dict = []
45 for m in arch:
46     relu_loss = np.zeros(10)
47     sigmoid_loss = np.zeros(10)
48     relu_grad = np.zeros(10)
49     sigmoid_grad = np.zeros(10)
50     arc_dict = {}
51     counter_for_list = 0
52
53     mlp_relu = MLPClassifier(hidden_layer_sizes=m,
54                               activation='relu',
55                               solver='sgd', max_iter=1,
56                               shuffle=True, learning_rate_init=0.01, momentum=0.0)
57
58     mlp_sigmoid = MLPClassifier(hidden_layer_sizes=m,
59                               activation='logistic',
60                               solver='sgd', max_iter=1,
61                               shuffle=True, learning_rate_init=0.01, momentum=0.0)
62
63     # Flatten input data
64     nsamples, nx, ny = X_train.shape
65     d2_train_dataset = X_train.reshape((nsamples, nx*ny))
66
67     nsamples, nx, ny = X_validate.shape
68     d2_validate_dataset = X_validate.reshape((nsamples, nx*ny))
69
70     nsamples, nx, ny = scaled_images_test.shape
71     d2_test_dataset = scaled_images_test.reshape((nsamples, nx*ny))
72
73     valid_accuracy = np.empty(shape=(10,))
74     training_accuracy = np.empty(shape=(10,))
75     Loss = np.empty(shape=(10,))
76
77     # For loop to 100 epochs
78     counter_10_epoch = 0
79     for i in range(1, 101):
80         # Shuffle training set after each epoch

```

```

77         random_perm = np.random.permutation(X_train.shape
    [0])
78         mini_batch_index = 0
79
80         # After the first epoch, save the previous
81         # weights in order to calculate magnitude of the gradient
82         # loss
83         if i != 1:
84             weights_relu_before = copy.deepcopy(mlp_relu.
85             coefs_)
86             weights_relu_before = weights_relu_before[0]
87             weights_sigmoid_before = copy.deepcopy(
88             mlp_sigmoid.coefs_)
89             weights_sigmoid_before =
90             weights_sigmoid_before[0]
91
92         # Until all data is fed to network, continue to
93         # feed batches.
94         while True:
95             indices = random_perm[mini_batch_index:
96             mini_batch_index + 500]
97             mlp_relu.partial_fit(d2_train_dataset[indices],
98             y_train[indices], np.unique(y_train))
99             mlp_sigmoid.partial_fit(d2_train_dataset[
100             indices], y_train[indices], np.unique(y_train))
101             mini_batch_index += 500
102
103             if mini_batch_index >= X_train.shape[0]:
104                 break
105
106             # Record the parameters for every 10 steps
107             if i % 10 == 0:
108                 weights_relu_after = mlp_relu.coefs_[0]
109                 weights_sigmoid_after = mlp_sigmoid.coefs_[0]
110                 relu_grad[counter_10_epoch] = np.linalg.norm(
111                     (weights_relu_before - weights_relu_after) * 100)
112                 sigmoid_grad[counter_10_epoch] = np.linalg.
113                 norm((weights_sigmoid_before - weights_sigmoid_after) *
114                     100)
115                 relu_loss[counter_10_epoch] = mlp_relu.loss_
116                 sigmoid_loss[counter_10_epoch] = mlp_sigmoid.
117                 loss_
118                 counter_10_epoch += 1

```

```
108     arc_dict['name'] = m
109     arc_dict['relu_loss_curve'] = relu_loss
110     arc_dict['sigmoid_loss_curve'] = sigmoid_loss
111     arc_dict['relu_grad_curve'] = relu_grad
112     arc_dict['sigmoid_grad_curve'] = sigmoid_grad
113     list_of_dict.append(arc_dict)
114
115 # Visualization Part
116 from utils import part3Plots
117 part3Plots(list_of_dict, save_dir='', filename='',
118             show_plot=True)
119
```

```
1 from sklearn.preprocessing import MinMaxScaler
2 from sklearn.model_selection import train_test_split
3 import matplotlib.pyplot as plt
4 from sklearn.neural_network import MLPClassifier
5 import numpy as np
6
7
8 train_images = np.load('train_images.npy')
9 train_labels = np.load('train_labels.npy')
10 test_images = np.load('test_images.npy')
11 test_labels = np.load('test_labels.npy')
12
13 train_images_reshaped = np.reshape(train_images, (30000,
14 28, 28))
14 test_images_reshaped = np.reshape(test_images, (
15 test_images.shape[0], 28, 28))
16
16 scaled_images = np.empty(shape=(30000, 28, 28))
17 scaled_images_test = np.empty(shape=(5000, 28, 28))
18 # In order to scale pixel values to [-1, 1], The
18 MaxAbsScaler was used.
19 counter = 0
20 scaler = MinMaxScaler(feature_range=(-1, 1))
21
22 for x in train_images_reshaped:
23     scaled_images[counter, :, :] = scaler.fit_transform(x)
24     counter += 1
25
26 counter = 0
27 for x in test_images_reshaped:
28     scaled_images_test[counter, :, :] = scaler.
28     fit_transform(x)
29     counter += 1
30
31 # In order to split 10 percent of the training data set to
31 validation
32 X_train, X_validate, y_train, y_validate =
32 train_test_split(scaled_images,
33
33     train_labels, test_size=0.1,
34
34     random_state=40, stratify=train_labels)
35 # parameter for the number of epoch
36 number_of_epoch = 20
37
```

```

38 # Chosen favourite architecture
39 arch_fav = (16, 128, 16)
40
41 average_loss_1 = np.zeros(number_of_epoch)
42 average_loss_01 = np.zeros(number_of_epoch)
43 average_loss_001 = np.zeros(number_of_epoch)
44
45 average_valid_scheduled = np.zeros(number_of_epoch)
46
47 average_valid_accuracy_1 = np.zeros(number_of_epoch)
48 average_valid_accuracy_01 = np.zeros(number_of_epoch)
49 average_valid_accuracy_001 = np.zeros(number_of_epoch)
50
51 counter_for_list = 0
52 arc_dict = {}
53
54 # There is no averaging in this part, so just one loop
      will be enough
55 for x in range(0, 1):
56     mlp_1 = MLPClassifier(hidden_layer_sizes=arch_fav,
                           activation='relu',
                           solver='sgd', max_iter=1, shuffle=
                           True, learning_rate_init=0.1, momentum=0.0)
57
58     mlp_01 = MLPClassifier(hidden_layer_sizes=arch_fav,
                           activation='relu',
                           solver='sgd', max_iter=1, shuffle=
                           True, learning_rate_init=0.01, momentum=0.0)
59
60     mlp_001 = MLPClassifier(hidden_layer_sizes=arch_fav,
                           activation='relu',
                           solver='sgd', max_iter=1, shuffle=
                           True, learning_rate_init=0.001, momentum=0.0)
61
62     mlp_scheduled = MLPClassifier(hidden_layer_sizes=
                           arch_fav, activation='relu',
                           solver='sgd', max_iter=1, shuffle=True
                           , learning_rate_init=0.1, momentum=0.0)
63
64     # Flatten input data
65     nsamples, nx, ny = X_train.shape
66     d2_train_dataset = X_train.reshape((nsamples, nx*ny))
67
68     nsamples, nx, ny = X_validate.shape
69     d2_validate_dataset = X_validate.reshape((nsamples, nx

```

```

73 *ny) )
74
75     nsamples, nx, ny = scaled_images_test.shape
76     d2_test_dataset = scaled_images_test.reshape( (
    nsamples, nx*ny) )
77
78     valid_accuracy_1 = np.empty(shape=(number_of_epoch,))
79     valid_accuracy_01 = np.empty(shape=(number_of_epoch,))
    )
80     valid_accuracy_001 = np.empty(shape=(number_of_epoch,
    ))
81
82     valid_accuracy_scheduled = np.empty(shape=((
    number_of_epoch))
83
84     Loss_1 = np.empty(shape=(number_of_epoch,))
85     Loss_01 = np.empty(shape=(number_of_epoch,))
86     Loss_001 = np.empty(shape=(number_of_epoch,))
87
88     # For loop to 200 epochs
89     counter_10_epoch = 0
90     for i in range(1, 201):
91         # Shuffle training set after each epoch
92         random_perm = np.random.permutation(X_train.shape
[0])
93         mini_batch_index = 0
94
95         # First change of learning rate for scheduled
learning method
96         if i == 70:
97             mlp_scheduled.set_params(learning_rate_init=0
.01)
98
99         # Second change of learning rate for scheduled
learning method
100        if i == 120:
101            mlp_scheduled.set_params(learning_rate_init=0
.001)
102
103        # Until all data is fed to network, continue to
fed batches.
104        while True:
105            indices = random_perm[mini_batch_index:
mini_batch_index + 500]
106            mlp_1.partial_fit(d2_train_dataset[indices],

```

File - C:\Users\berkenutku\Desktop\496_hw\MLP_part4.py

```
138
139 arc_dict['name'] = 'arch_3'
140 arc_dict['loss_curve_1'] = average_loss_1
141 arc_dict['loss_curve_01'] = average_loss_01
142 arc_dict['loss_curve_001'] = average_loss_001
143 arc_dict['val_acc_curve_1'] = average_valid_accuracy_1
144 arc_dict['val_acc_curve_01'] = average_valid_accuracy_01
145 arc_dict['val_acc_curve_001'] = average_vali
    d_accuracy_001
146
147 # Visualization Part
148 from utils import part4Plots
149 part4Plots(arc_dict, save_dir='', filename='', show_plot=
    True)
150
151 plt.plot(average_valid_scheduled)
152 plt.ylabel('Validation Accuracy')
153 plt.show()
154
155 print(mlp_scheduled.score(d2_train_dataset, y_train))
156
```